

Chroma self-querying

Chroma is a database for building AI applications with embeddings.

In the notebook we'll demo the `SelfQueryRetriever` wrapped around a Chroma vector store.

Creating a Chroma vectorstore

First we'll want to create a Chroma VectorStore and seed it with some data. We've created a small demo set of documents that contain summaries of movies.

NOTE: The self-query retriever requires you to have `lark` installed (`pip install lark`). We also need the `chromadb` package.

```
#!pip install lark
```



```
#!pip install chromadb
```

We want to use `OpenAIEmbeddings` so we have to get the OpenAI API Key.

```
import os  
import getpass
```

```
os.environ["OPENAI_API_KEY"] = getpass.getpass("OpenAI API Key:")
```

OpenAI API Key:

```
from langchain.schema import Document
from langchain.embeddings.openai import OpenAIEmbeddings
from langchain.vectorstores import Chroma

embeddings = OpenAIEmbeddings()
```

```
docs = [
    Document(
        page_content="A bunch of scientists bring back dinosaurs and mayhem breaks loose",
        metadata={"year": 1993, "rating": 7.7, "genre": "science fiction"},
    ),
    Document(
        page_content="Leo DiCaprio gets lost in a dream within a dream within a dream within a ...",
        metadata={"year": 2010, "director": "Christopher Nolan", "rating": 8.2},
    ),
    Document(
        page_content="A psychologist / detective gets lost in a series of dreams within dreams within dreams and Inception reused the idea",
        metadata={"year": 2006, "director": "Satoshi Kon", "rating": 8.6},
    ),
    Document(
        page_content="A bunch of normal-sized women are supremely wholesome and some men pine after them",
        metadata={"year": 2019, "director": "Greta Gerwig", "rating": 8.3},
    ),
]
```

```
Document(  
    page_content="Toys come alive and have a blast doing so",  
    metadata={"year": 1995, "genre": "animated"},  
)  
Document(  
    page_content="Three men walk into the Zone, three men walk out of the Zone",  
    metadata={  
        "year": 1979,  
        "rating": 9.9,  
        "director": "Andrei Tarkovsky",  
        "genre": "science fiction",  
        "rating": 9.9,  
    },  
)  
]  
vectorstore = Chroma.from_documents(docs, embeddings)
```

Using embedded DuckDB without persistence: data will be transient

Creating our self-querying retriever

Now we can instantiate our retriever. To do this we'll need to provide some information upfront about the metadata fields that our documents support and a short description of the document contents.

```
from langchain.llms import OpenAI  
from langchain.retrievers.self_query.base import SelfQueryRetriever  
from langchain.chains.query_constructor.base import AttributeInfo
```

```
metadata_field_info = [  
    AttributeInfo(  
        name="genre",  
        description="The genre of the movie",  
        type="string or list[string]",  
    ),  
    AttributeInfo(  
        name="year",  
        description="The year the movie was released",  
        type="integer",  
    ),  
    AttributeInfo(  
        name="director",  
        description="The name of the movie director",  
        type="string",  
    ),  
    AttributeInfo(  
        name="rating", description="A 1-10 rating for the movie", type="float"  
    ),  
]  
document_content_description = "Brief summary of a movie"  
llm = OpenAI(temperature=0)  
retriever = SelfQueryRetriever.from_llm(  
    llm, vectorstore, document_content_description, metadata_field_info, verbose=True  
)
```

Testing it out

And now we can try actually using our retriever!

```
# This example only specifies a relevant query
retriever.get_relevant_documents("What are some movies about dinosaurs")
```

```
query='dinosaur' filter=None
```

```
[Document(page_content='A bunch of scientists bring back dinosaurs and mayhem breaks loose', metadata=
{'year': 1993, 'rating': 7.7, 'genre': 'science fiction'}),
 Document(page_content='Toys come alive and have a blast doing so', metadata={'year': 1995, 'genre':
'animated'}),
 Document(page_content='A psychologist / detective gets lost in a series of dreams within dreams within
dreams and Inception reused the idea', metadata={'year': 2006, 'director': 'Satoshi Kon', 'rating': 8.6}),
 Document(page_content='Leo DiCaprio gets lost in a dream within a dream within a dream within a ...',
metadata={'year': 2010, 'director': 'Christopher Nolan', 'rating': 8.2})]
```

```
# This example only specifies a filter
retriever.get_relevant_documents("I want to watch a movie rated higher than 8.5")
```

```
query=' ' filter=Comparison(comparator=Comparator.GT: 'gt', attribute='rating', value=8.5)
```

```
[Document(page_content='A psychologist / detective gets lost in a series of dreams within dreams within
```

```
dreams and Inception reused the idea', metadata={'year': 2006, 'director': 'Satoshi Kon', 'rating': 8.6}),  
    Document(page_content='Three men walk into the Zone, three men walk out of the Zone', metadata={'year':  
1979, 'rating': 9.9, 'director': 'Andrei Tarkovsky', 'genre': 'science fiction'})]]
```

This example specifies a query and a filter

```
retriever.get_relevant_documents("Has Greta Gerwig directed any movies about women")
```

```
query='women' filter=Comparison(comparator=<Comparator.EQ: 'eq'>, attribute='director', value='Greta  
Gerwig')
```

```
[Document(page_content='A bunch of normal-sized women are supremely wholesome and some men pine after  
them', metadata={'year': 2019, 'director': 'Greta Gerwig', 'rating': 8.3})]
```

This example specifies a composite filter

```
retriever.get_relevant_documents(  
    "What's a highly rated (above 8.5) science fiction film?"  
)
```

```
query=' ' filter=Operation(operator=<Operator.AND: 'and'>, arguments=[Comparison(comparator=  
<Comparator.EQ: 'eq'>, attribute='genre', value='science fiction'), Comparison(comparator=<Comparator.GT:  
'gt'>, attribute='rating', value=8.5)])
```

```
[Document(page_content='Three men walk into the Zone, three men walk out of the Zone', metadata={'year': 1979, 'rating': 9.9, 'director': 'Andrei Tarkovsky', 'genre': 'science fiction'})]
```

This example specifies a query and composite filter

```
retriever.get_relevant_documents(  
    "What's a movie after 1990 but before 2005 that's all about toys, and preferably is animated"  
)
```

```
query='toys' filter=Operation(operator=<Operator.AND: 'and'>, arguments=[Comparison(comparator=  
<Comparator.GT: 'gt'>, attribute='year', value=1990), Comparison(comparator=<Comparator.LT: 'lt'>,   
attribute='year', value=2005), Comparison(comparator=<Comparator.EQ: 'eq'>, attribute='genre',   
value='animated')])
```

```
[Document(page_content='Toys come alive and have a blast doing so', metadata={'year': 1995, 'genre':   
'animated'})]
```

Filter k

We can also use the self query retriever to specify `k`: the number of documents to fetch.

We can do this by passing `enable_limit=True` to the constructor.

```
retriever = SelfQueryRetriever.from_llm(  
    llm,  
    vectorstore,  
    document_content_description,  
    metadata_field_info,  
    enable_limit=True,  
    verbose=True,  
)
```

```
# This example only specifies a relevant query  
retriever.get_relevant_documents("what are two movies about dinosaurs")
```

```
query='dinosaur' filter=None
```

```
[Document(page_content='A bunch of scientists bring back dinosaurs and mayhem breaks loose', metadata=  
{ 'year': 1993, 'rating': 7.7, 'genre': 'science fiction'}),  
 Document(page_content='Toys come alive and have a blast doing so', metadata={ 'year': 1995, 'genre':  
'animated'}),  
 Document(page_content='A psychologist / detective gets lost in a series of dreams within dreams within  
dreams and Inception reused the idea', metadata={ 'year': 2006, 'director': 'Satoshi Kon', 'rating': 8.6}),  
 Document(page_content='Leo DiCaprio gets lost in a dream within a dream within a dream within a ...',  
metadata={ 'year': 2010, 'director': 'Christopher Nolan', 'rating': 8.2})]
```