

JSONFormer

JSONFormer is a library that wraps local HuggingFace pipeline models for structured decoding of a subset of the JSON Schema.

It works by filling in the structure tokens and then sampling the content tokens from the model.

Warning - this module is still experimental

```
pip install --upgrade jsonformer > /dev/null
```

HuggingFace Baseline

First, let's establish a qualitative baseline by checking the output of the model without structured decoding.

```
import logging

logging.basicConfig(level=logging.ERROR)
```

```
from typing import Optional
from langchain.tools import tool
import os
import json
import requests
```

```
HF_TOKEN = os.environ.get("HUGGINGFACE_API_KEY")
```

```
@tool
```

```
def ask_star_coder(query: str, temperature: float = 1.0, max_new_tokens: float = 250):  
    """Query the BigCode StarCoder model about coding questions."""  
    url = "https://api-inference.huggingface.co/models/bigcode/starcoder"  
    headers = {  
        "Authorization": f"Bearer {HF_TOKEN}",  
        "content-type": "application/json",  
    }  
    payload = {  
        "inputs": f"{query}\n\nAnswer:",  
        "temperature": temperature,  
        "max_new_tokens": int(max_new_tokens),  
    }  
    response = requests.post(url, headers=headers, data=json.dumps(payload))  
    response.raise_for_status()  
    return json.loads(response.content.decode("utf-8"))
```

```
prompt = """You must respond using JSON format, with a single action and single action input.  
You may 'ask_star_coder' for help on coding problems.
```

```
{arg_schema}
```

```
EXAMPLES
```

```
----
```

```
Human: "So what's all this about a GIL?"
```

```
AI Assistant:{{
```

```
    "action": "ask_star_coder",  
    "action_input": {"query": "What is a GIL?", "temperature": 0.0, "max_new_tokens": 100}}  
}}
```

Observation: "The GIL is python's Global Interpreter Lock"

Human: "Could you please write a calculator program in LISP?"

```
AI Assistant:{{
    "action": "ask_star_coder",
    "action_input": {"query": "Write a calculator program in LISP", "temperature": 0.0, "max_new_tokens":
250}}
}}
```

Observation: "(defun add (x y) (+ x y))\n(defun sub (x y) (- x y))"

Human: "What's the difference between an SVM and an LLM?"

```
AI Assistant:{{
    "action": "ask_star_coder",
    "action_input": {"query": "What's the difference between SGD and an SVM?", "temperature": 1.0,
"max_new_tokens": 250}}
}}
```

Observation: "SGD stands for stochastic gradient descent, while an SVM is a Support Vector Machine."

BEGIN! Answer the Human's question as best as you are able.

Human: 'What's the difference between an iterator and an iterable?'

```
AI Assistant:"".format(
    arg_schema=ask_star_coder.args
)
```

```
from transformers import pipeline
from langchain.llms import HuggingFacePipeline

hf_model = pipeline(
    "text-generation", model="cerebras/Cerebras-GPT-590M", max_new_tokens=200
)

original_model = HuggingFacePipeline(pipeline=hf_model)
```

```
generated = original_model.predict(prompt, stop=["Observation:", "Human:"])  
print(generated)
```

```
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
```

```
'What's the difference between an iterator and an iterable?'
```

That's not so impressive, is it? It didn't follow the JSON format at all! Let's try with the structured decoder.

JSONFormer LLM Wrapper

Let's try that again, now providing the Action input's JSON Schema to the model.

```
decoder_schema = {  
    "title": "Decoding Schema",  
    "type": "object",  
    "properties": {  
        "action": {"type": "string", "default": ask_star_coder.name},  
        "action_input": {  
            "type": "object",  
            "properties": ask_star_coder.args,  
        },  
    },  
}
```

```
from langchain.experimental.llms import JsonFormer

json_former = JsonFormer(json_schema=decoder_schema, pipeline=hf_model)
```

```
results = json_former.predict(prompt, stop=["Observation:", "Human:"])
print(results)
```

```
{
  "action": "ask_star_coder",
  "action_input": {
    "query": "What's the difference between an iterator and an iter",
    "temperature": 0.0,
    "max_new_tokens": 50.0
  }
}
```

Voila! Free of parsing errors.